

---

# Process Types

- Sometimes OS may swap a blocked process to disk to free up more memory
- Or to improve process mix
- This is called **Swapping**

---

# Reference

- 10.3.2 Process Management System Calls in UNIX
- Modern Operating System
  - 2<sup>nd</sup> Edition, Andrew S. Tanenbaum

---

# Process Creation

- Includes
  - Build kernel data structures
  - Allocate memory
- Reasons to create a process
  - Submit a new batch job/Start program
  - User logs on to the system
  - OS creates on behalf of a user (printing)
  - Spawned by existing process

# Process Termination

- Batch job issues Halt instruction
- User logs off
- Process executes a service request to terminate
- Parent terminates so child processes terminate
- Operating system intervention
  - such as when deadlock occurs
- Error and fault conditions
  - E.g. memory unavailable, protection error, arithmetic error, I/O failure, invalid instruction

---

# Unix Process Creation

- When the system starts up it is running in kernel mode
- There is only one process, the initial process.
- At the end of system initialization, the initial process starts up another kernel process.
- The *init* kernel process has a process identifier of 1.

---

# Process Creation

- These new processes may themselves go on to create new processes.
- All of the processes in the system are descended from the *init* kernel thread.
- You can see the family relationship between the running processes in a Linux system using the *ps tree* command
- A new process is created by a **fork()** system call

---

# Compiling C ++ code

- `g++ test.cpp -o Output`
- Running the code:
- `./Output`

---

# The `fork()` system call

At the end of the system call there is a new process waiting to run once the scheduler chooses it

- A new data structure is allocated
- The new process is called the child process.
- The existing process is called the parent process.
- The parent gets the child's pid returned to it.
- The child gets 0 returned to it.
- Both parent and child execute at the same point after `fork()` returns



---

# Unix Process Control

```
int main()  
{  
    int pid;  
    int x = 0;  
  
    x = x + 1;  
    pid = fork();  
    x = 3;  
    printf("%d",x);  
}
```

# But we want the child process to do something else...

```
int pid;  
int status = 0;  
  
if (pid = fork()) {  
    /* parent */  
    .....  
    pid = wait(&status);  
} else {  
    /* child */  
    .....  
    exit(status);  
}
```

*The **fork** syscall returns a zero to the child and the child process ID to the parent*

*Parent uses **wait** to*

*sleep until child returns and status.*  
**Fork** creates an exact copy of the parent process

*Wait variants allow wait on a specific pid.*  
**Wait** variants allow wait on a specific pid.

*Child process passes status back to parent on **exit**, to report success/failure*

---

# Child Process Inherits

- Stack
- Memory
- Environment
- Open file descriptors.
- Current working directory
- Resource limits
- Root directory

---

# Child process DOESNOT Inherit

- Process ID
- Different parent process ID
- Process times
- Own copy of file descriptors
- Resource utilization (initialized to zero)

---

## How can a parent and child process communicate?

- Through any of the normal IPC mechanism schemes.
- But have special ways to communicate
- For example
  - The variables are replicas
  - The parent receives the exit status of the child

---

# The wait() System Call

- **A child program returns a value to the parent, so the parent must arrange to receive that value**
- **The wait() system call serves this purpose**
  - ❑ `pid_t wait(int *status)`
  - ❑ it puts the parent to sleep waiting for a child's result
  - ❑ when a child calls `exit()`, the OS unblocks the parent and returns the value passed by `exit()` as a result of the wait call (along with the pid of the child)
  - ❑ if there are no children alive, `wait()` returns immediately
  - ❑ also, if there are zombies, `wait()` returns one of the values immediately (and deallocates the zombie)

---

# What is a zombie?

- In the interval between the child terminating and the parent calling wait(), the child is said to be a 'zombie'.
- Even though its not running its taking up an entry in the process table.
- The process table has a limited number of entries.

---

# What is a zombie?

- If the parent terminates without calling `wait()`, the child is adopted by `init`.

The solution is:

- Ensure that your parent process calls `wait()` or `waitpid` or etc, for every child process that terminates.



# exit()

## **void exit(int *status*);**

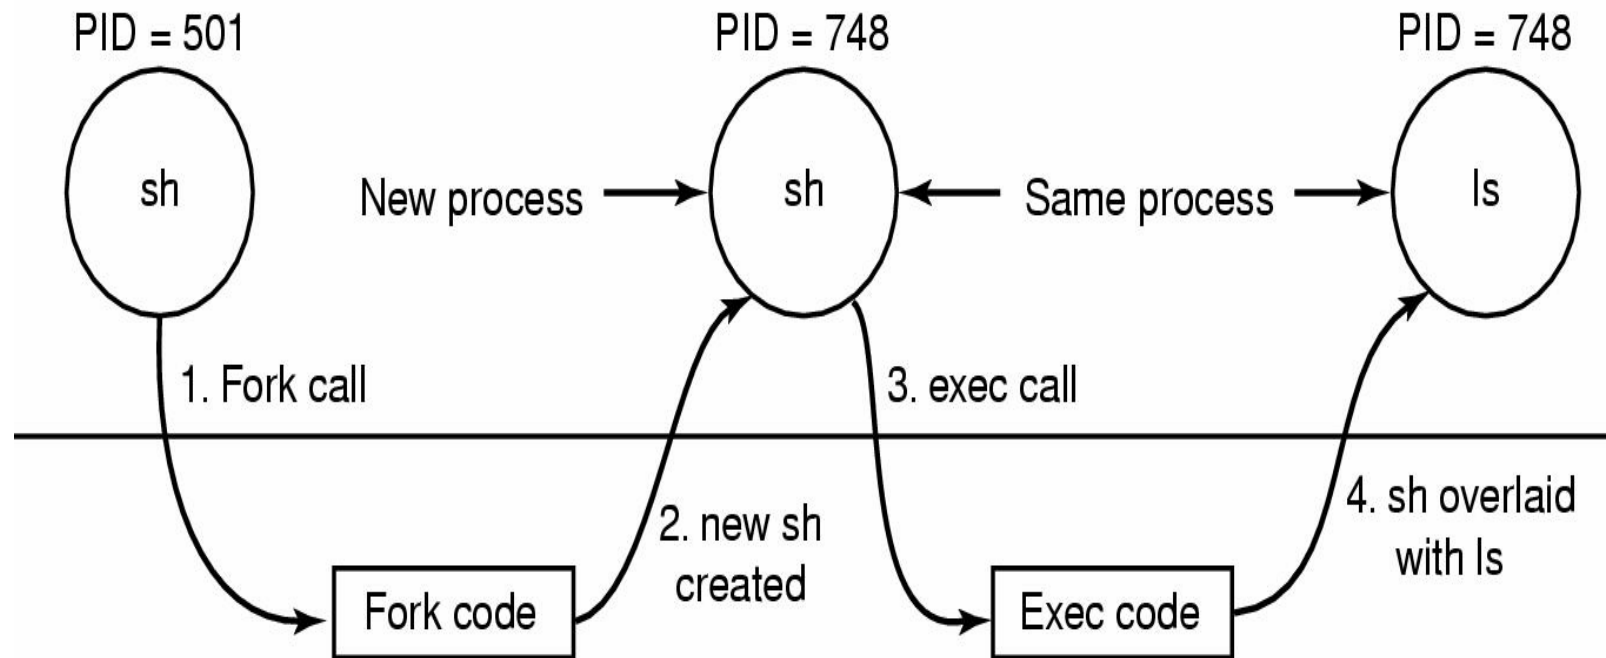
- After the program finishes execution, it calls *exit()*
- This system call:
  - takes the “result” of the program as an argument
  - closes all open files, connections, etc.
  - deallocates memory
  - deallocates most of the OS structures supporting the process
  - checks if parent is alive:
    - If so, it holds the result value until parent requests it, process does not really die, but it enters the zombie/defunct state
    - If not, it deallocates all data structures, the process is dead

---

## execv()

- We usually want the child process to run some other executable
- For Example, *ls*

## The `ls` Command



Steps in executing the command `ls` type to the shell

---

## execv

- `int execv(const char *path, char *const argv[]);`
- the current process image with a new process image.
- ***path*** is the filename to be executed by the child process
- When a C-language program is executed as a result of this call, it is entered as a C-language function call as follows:
  - `int main (int argc, char *argv[]);`
- The ***argv*** array is terminated by a null pointer.
- The null pointer terminating the ***argv*** array is not counted in ***argc***.

---

```
while (TRUE) {                                /* repeat forever */
    type_prompt( );                            /* display prompt on the screen */
    read_command(command, params);             /* read input line from keyboard */

    pid = fork( );                            /* fork off a child process */
    if (pid < 0) {
        printf("Unable to fork");             /* error condition */
        continue;                             /* repeat the loop */
    }

    if (pid != 0) {
        waitpid (-1, &status, 0);             /* parent waits for child */
    } else {
        execve(command, params, 0);           /* child does the work */
    }
}
```